
Specification of Precision of Floating-point and Integer Types

Paul A. Bristow
Christopher Kormanyos

Copyright © 2013 Paul A. Bristow, Christopher Kormanyos

Distributed under the Boost Software License, Version 1.0. (See accompanying file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

Table of Contents

Abstract 2

Background 3

Introduction 4

How to specify extended precision constants - Q? 5

Specifying Precision 6

 Existing Fixed precision integer types 7

 Proposed new section 7

References 8

Version Info 9

ISO/IEC JTC1 SC22 WG21 N??? - 2013-4-??



Important

This is NOT an official Boost library.



Note

Comments and suggestions to Paul.A.Bristow pbristow@hetp.u-net.com.

Abstract

It is proposed to add several optional typedefs with fixed precisions for floating-point types including `float32_t`, `float64_t`, `_float128_t` (similar to `int64_t` for integer types).

These will be defined in the global and `std` namespaces.

And also to provide additional suffix(es) to specify extended precision constants to suit precisions higher than `long double`.

The objectives are to makes it easier to use higher-precision, to reduce errors in precision, and to improve portability.

Background

Mathematical functions were added to the C++11 libraries via technical report TR1; it is now proposed to fix these into the next C++1Y standard.¹

Other mathematical special functions are also now proposed, for example, [A proposal to add special mathematical functions according to the ISO/IEC 80000-2:2009 standard Document number: N3494 Version: 1.0 Date: 2012-12-19](#)

The [Boost.Math](#) library was accepted into [Boost](#) several years ago. It implements many of the functions in both documents mentioned above and has become quite widely used.

With the acceptance and release of [Boost.Multiprecision](#) that provides much higher precision than built-in `long double` with `cpp_dec_float` employing a variety of backends including the well-established [GNU Multiple Precision Arithmetic Library](#) and [GNU MPFR library](#) libraries as well as a full open-license backend developed from the `e_float` (TOMS Algorithm 910) library by Christopher Kormanyos and John Maddock.

Since [Boost.Multiprecision](#) and [Boost.Math](#) work seamlessly, allowing a `float_type` typedef to be switched from a built-in type to hundreds of decimal digits; then all the special functions and distributions can be used at any chosen precision.

Other users and domains are finding the need and utility of [decimal](#) and [binary fixed-point](#).

Of course, moving away from hardware supported types to software using templates carries a small price at compile-time, and a much bigger price at runtime.

All these development have made C++ much more attractive to the scientific and engineering community, especially those needing higher (or lower) precision for some (if not all) of the calculations, previously the domain covered by [Wolfram Mathematica](#), [MATLAB](#) and others where the precision can be arbitrarily chosen.

¹ Conditionally-supported Special Math Functions for C++14, N3584, Walter E. Brown

Introduction

These developments have also revealed a need for more standard ways to specify precision, especially for extended precision, and to improve portability and give more confidence that the actual precision can be predicted. For example, support has been expressed on the [Boost list discussion of precise floating-point types](#)

The reasons for this need are analogous to those that led to the introduction of fixed integer size like `int64_t`.

Recent specification of fixed-size integer types in C99, C11 and C++11 and [C++ draft specification](#) has drastically improved integer algorithm portability and range.

Similar specification of fixed-size floating-point types could potentially improve the C++ language significantly, especially in the scientific and engineering communities.

One example of how fixed-size integer types have proved invaluable is described by Robert Ramey [Usefulness of fixed integer sizes in portability \(for Boost serialization library\)](#).

“Fundamental types in C++ are `unsigned char`, `signed char`, `unsigned short int`, `signed short int`, ... `unsigned long`, `signed long`. In addition to the above some compilers define `int32_t`, and other as fundamental types. It is a unfortunate accident of history that the nomenclature is confusing. It is an unfortunate original design choice that this size of `int`, `char` etc were not defined as a specific number of bits. However at the time there were in common usage machines with 9, 16, 18, 24, 32, 36 and 48 bit words. What else were the authors to do? It is common among programers to define types `int16_t`, ..., etc using the `typedef` facility to map integers of a specific size between machines. This does no harm and can facilitate portability. However it in no way alters the fundamental types that are available on a given platform.”

How to specify extended precision constants - Q?

Recent discussion on extended precision floating-point types in C++ has also raised the issue of how to specify constant values with a precision greater than `long double`, now signified by the suffix `L`.

One obvious way is to add `Q` or `q` suffixes to signify that a constant has at least 128-bits (about 40 decimal digits) of precision.

There may also be a need for 256-bit (about 80 decimal digits) precision, and perhaps 512-bits (about 155 decimal digits) precision.

At present, the only way to provide constant values is to use a string to extended-precision type conversion.

This `from_string` method is used for [Boost.Math](#), [Boost.Multiprecision](#) and [GCC libquadmath](#), for example.

It would also be useful to have a method of interrogating the size of types, similar to that provided by [GCC 3.7.2 Common Predefined Macros](#), for example, `__SIZEOF_LONG_DOUBLE__` (but is not defined for `__float128` nor `__float80`)

We refer to floating-point types with fixed precision such as 24, 54, 113 or more binary significand digits, (and possibly even extending beyond these to potential multiprecision types).

Specifying Precision

One could envision two ways to name the fixed-precision types:

- `float24_t`, `float53_t`, `float113_t`, ...
- `float32_t`, `float64_t`, `float128_t`, ...

The first set above is intuitively coined from IEE754:2008. It is also consistent with the gist of `std::uint32_t`, et al in so far as the number of binary digits of *significand* precision is contained within the name of the data type.

On the other hand, the second set using the size of the *whole type* may seem more intuitive to users. The exact layout and number of significand and exponent bits can be confirmed by checking `std::numeric_limits<type>::is_iec559 == true`.

With the availability of Boost.Multiprecision, C++ programmers can now easily switch to using floating-point types that give far more decimal digits of precision (hundreds) than the built-in types `float`, `double` and `long double`.

And portability is also reduced. For example, suppose we wish to achieve a precision higher than the most common IEEE 64-bit floating-point type supported by the X86 chipsets normally used for `double`. http://en.wikipedia.org/wiki/Double_precision providing a precision of between 15 to 17 decimal digits.

The options for `long double` are many.

At least one popular compiler treats `long double` exactly as `double` (as permitted by the C++ Standard which does not prescribe the precision for any floating-point (or integer) types, leaving them to be implementation-defined).

However the [Intel X8087 chipset](#) does do calculations using internal 80-bit registers, increasing the significand from 53 to 63 bits, and gaining about 3 decimal digits precision from 18 and 21.

Some hardware, for example [Sparc](#), provides a 128-bit quadruple precision floating-point chip.

As of gcc 4.3, a quadruple precision is also supported on x86, but as the nonstandard type `__float128` rather than `long double`.

[Darwin](#) `long double` uses a double-double format developed first by [Keith Briggs](#). This gives about 106-bits of precision (about 33 decimal digits) but has rather odd behaviour at the extremes making implementation of `std::numeric_limits<>::epsilon()` problematic.

Clang uses a similar technique

```
#ifdef __clang__
    typedef struct { long double x, y; } __float128;
#endif
```

as described in [Clang float128](#).

If we wish to ensure that we use all 80 bits to calculate [Extended precision](#) we would use a `typedef float80_t`.

If the compiler could not generate code this type directly, then it would substitute software emulation, perhaps using a Boost.Multiprecision type `cpp_dec_float_21`.

Similarly if a quadrupole precision of 16-byte 128-bit [format](#) is desired, the specification of `float128_t` will either direct the compiler to generate code using the hardware, or it will do this using software emulation. This might be generated by the compiler for GCC or delegated to a `cpp_bin_float_128` type (under development for [Boost.Multiprecision](#)).

GNU C supports additional floating types, `__float80` and `__float128` to support 80-bit (XFmode) and 128-bit (TFmode) floating types.

Existing Fixed precision integer types

18.4 Integer types [cstdint]

18.4.1 Header <stdint> synopsis [cstdint.syn]

```
namespace std
{
    typedef signed integer type int8_t; // optional
    typedef signed integer type int16_t; // optional
    typedef signed integer type int32_t; // optional
    typedef signed integer type int64_t; // optional
}
```

Proposed new section

Add the following text to <stdint>



Note

It is not obvious where these typedef should reside. The obvious place is <stdint> but int implies integer types. (or <stdfloat>?)

18.4? Arithmetic types [stdfloat] (or cstdarith) 18.4.2? Header <stdfloat> synopsis [stdfloat.syn]

```
namespace std {
    typedef signed integer type float_32_t; // optional
    typedef signed integer type float_64_t; // optional
    typedef signed integer type float_80_t; // optional
    typedef signed integer type float_128_t; // optional
    typedef signed integer type float_256_t; // optional
} // namespace std
```



Note

Others might also be defined here?

It is not proposed to make any change to `std::numeric_limits`.

It is obviously highly desirable that `std::numeric_limits` is specialized for all floating or fixed-point types. And experience with [Boost.Math](#) and [Boost.Multiprecision](#) is that the normal set of trig and others useful functions is also essential to make the type useful in real-life.

Programs can then use this to determine if a floating-point type is IEEE 754 using `std::numeric_limits<>::is_iec559`.

References

isocpp.org C++ papers and mailings

[C++ Binary Fixed-Point Arithmetic, N3352, Lawrence Crowl](#)

[Proposal to Add Decimal Floating Point Support to C++, N3407 Dietmar Kuhl](#)

The C committee is working on a Decimal TR as TR 24732. The decimal support in C uses built-in types `_Decimal32`, `_Decimal64`, and `_Decimal128`. [128-bit decimal floating point in IEEE 754:2008](#)

[lists binary16, 32, 64 and 128](#)

(and also decimal 32, 64, and 128) [IEEE Std 754-2008](#)

[IEEE Standard for Floating-point Arithmetic, IEEE Std 754-2008](#)

[How to Read Floating Point Numbers Accurately, William D Clinger](#)

[Conditionally-supported Special Math Functions for C++14, N3584, Walter E. Brown](#)

[Walter E. Brown, Opaque Typedefs](#)

[Specification of Extended Precision Floating-point and Integer Types, Christopher Kormanyos, John Maddock](#)

[X8087 notes](#)

Version Info

Last edit to Quickbook file precision.qbk was at 02:51:17 PM on 2013-Mar-21.



Tip

This should appear on the pdf version (but may be redundant on a html version where the last edit date is on the first (home) page).



Warning

Home page "Last revised" is GMT, not local time. Last edit date is local time.